

HW2

OmerFarukYahsi

17.4 Time spans

We'll learn about how arithmetic with dates works, including subtraction, addition, and division. Along the way, you'll learn about three important classes that represent time spans:

- **Durations**, which represent an exact number of seconds.
- **Periods**, which represent human units like weeks and months.
- **Intervals**, which represent a starting and ending point.

How do you pick between duration, periods, and intervals? As always, pick the simplest data structure that solves your problem. If you only care about physical time, use a duration; if you need to add human times, use a period; if you need to figure out how long a span is in human units, use an interval.

This chapter will focus on the **lubridate** package, which makes it easier to work with dates and times in R. As of the latest tidyverse release, lubridate is part of core tidyverse. We will also need nycflights13 for practice data.

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.3      v readr      2.1.4
v forcats    1.0.0      v stringr    1.5.0
v ggplot2    3.4.3      v tibble     3.2.1
v lubridate  1.9.3      v tidyr      1.3.0
v purrr      1.0.2

-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
library(nycflights13)
library(lubridate)
```

17.4.1 Durations

In R, when you subtract two dates, you get a `difftime` object:

```
# How old is Hadley?
h_age <- today() - ymd("1979-10-14")

h_age
```

Time difference of 16108 days

A `difftime` class object records a time span of seconds, minutes, hours, days, or weeks. This ambiguity can make `difftimes` a little painful to work with, so `lubridate` provides an alternative which always uses seconds: the **duration**.

```
as.duration(h_age)
```

```
[1] "1391731200s (~44.1 years)"
```

Durations come with a bunch of convenient constructors:

```
dseconds(15)
```

```
[1] "15s"
```

```
dminutes(10)
```

```
[1] "600s (~10 minutes)"
```

```
dhours(c(12, 24))
```

```
[1] "43200s (~12 hours)" "86400s (~1 days)"
```

```
ddays(0:5)
```

```
[1] "0s" "86400s (~1 days)" "172800s (~2 days)"  
[4] "259200s (~3 days)" "345600s (~4 days)" "432000s (~5 days)"
```

```
dweeks(3)
```

```
[1] "1814400s (~3 weeks)"
```

```
dyears(1)
```

```
[1] "31557600s (~1 years)"
```

Durations always record the time span in seconds. Larger units are created by converting minutes, hours, days, weeks, and years to seconds: 60 seconds in a minute, 60 minutes in an hour, 24 hours in a day, and 7 days in a week. Larger time units are more problematic. A year uses the "average" number of days in a year, i.e. 365.25. There's no way to convert a month to a duration, because there's just too much variation.

You can add and multiply durations:

```
2 * dyears(1)
```

```
[1] "63115200s (~2 years)"
```

```
dyears(1) + dweeks(12) + dhours(15)
```

```
[1] "38869200s (~1.23 years)"
```

You can add and subtract durations to and from days:

```
tomorrow <- today() + ddays(1)
```

```
tomorrow
```

```
[1] "2023-11-21"
```

```
last_year <- today() - dyears(1)
```

```
last_year
```

```
[1] "2022-11-19 18:00:00 UTC"
```

However, because durations represent an exact number of seconds, sometimes you might get an unexpected result:

```
one_am <- ymd_hms("2026-03-08 01:00:00", tz = "America/New_York")
```

```
one_am
```

```
[1] "2026-03-08 01:00:00 EST"
```

```
one_am + ddays(1)
```

```
[1] "2026-03-09 02:00:00 EDT"
```

Why is one day after 1am March 8, 2am March 9? If you look carefully at the date you might also notice that the time zones have changed. March 8 only has 23 hours because it's when DST starts, so if we add a full days worth of seconds we end up with a different time.

17.4.2 Periods

To solve this problem, lubridate provides **periods**. Periods are time spans but don't have a fixed length in seconds, instead they work with "human" times, like days and months. That allows them to work in a more intuitive way:

```
one_am
```

```
[1] "2026-03-08 01:00:00 EST"
```

```
one_am + days(1)
```

```
[1] "2026-03-09 01:00:00 EDT"
```

Like durations, periods can be created with a number of friendly constructor functions.

```
hours(c(12, 24))
```

```
[1] "12H 0M 0S" "24H 0M 0S"
```

```
days(7)
```

```
[1] "7d 0H 0M 0S"
```

```
months(1:6)
```

```
[1] "1m 0d 0H 0M 0S" "2m 0d 0H 0M 0S" "3m 0d 0H 0M 0S" "4m 0d 0H 0M 0S"  
[5] "5m 0d 0H 0M 0S" "6m 0d 0H 0M 0S"
```

You can add and multiply periods:

```
10 * (months(6) + days(1))
```

```
[1] "60m 10d 0H 0M 0S"
```

```
days(50) + hours(25) + minutes(2)
```

```
[1] "50d 25H 2M 0S"
```

And of course, add them to dates. Compared to durations, periods are more likely to do what you expect:

```
# A leap year  
ymd("2024-01-01") + dyears(1)
```

```
[1] "2024-12-31 06:00:00 UTC"
```

```
ymd("2024-01-01") + years(1)
```

```
[1] "2025-01-01"
```

```
# Daylight saving time
one_am + ddays(1)
```

```
[1] "2026-03-09 02:00:00 EDT"
```

```
one_am + days(1)
```

```
[1] "2026-03-09 01:00:00 EDT"
```

The times are represented in a slightly odd format, so we use modulus arithmetic to pull out the hour and minute components. Once we've created the date-time variables, we focus in on the variables we'll explore in the rest of the chapter.

```
make_datetime_100 <- function(year, month, day, time) {
  make_datetime(year, month, day, time %/% 100, time %% 100)
}

flights_dt <- flights |>
  filter(!is.na(dep_time), !is.na(arr_time)) |>
  mutate(
    dep_time = make_datetime_100(year, month, day, dep_time),
    arr_time = make_datetime_100(year, month, day, arr_time),
    sched_dep_time = make_datetime_100(year, month, day, sched_dep_time),
    sched_arr_time = make_datetime_100(year, month, day, sched_arr_time)
  ) |>
  select(origin, dest, ends_with("delay"), ends_with("time"))

flights_dt
```

```
# A tibble: 328,063 x 9
```

	origin	dest	dep_delay	arr_delay	dep_time	sched_dep_time
	<chr>	<chr>	<dbl>	<dbl>	<dtm>	<dtm>
1	EWR	IAH	2	11	2013-01-01 05:17:00	2013-01-01 05:15:00
2	LGA	IAH	4	20	2013-01-01 05:33:00	2013-01-01 05:29:00
3	JFK	MIA	2	33	2013-01-01 05:42:00	2013-01-01 05:40:00
4	JFK	BQN	-1	-18	2013-01-01 05:44:00	2013-01-01 05:45:00
5	LGA	ATL	-6	-25	2013-01-01 05:54:00	2013-01-01 06:00:00
6	EWR	ORD	-4	12	2013-01-01 05:54:00	2013-01-01 05:58:00
7	EWR	FLL	-5	19	2013-01-01 05:55:00	2013-01-01 06:00:00

```

8 LGA      IAD          -3      -14 2013-01-01 05:57:00 2013-01-01 06:00:00
9 JFK      MCO          -3      -8 2013-01-01 05:57:00 2013-01-01 06:00:00
10 LGA     ORD          -2        8 2013-01-01 05:58:00 2013-01-01 06:00:00
# i 328,053 more rows
# i 3 more variables: arr_time <dtm>, sched_arr_time <dtm>, air_time <dbl>

```

Let's use periods to fix an oddity related to our flight dates. Some planes appear to have arrived at their destination *before* they departed from New York City.

```

flights_dt |>
  filter(arr_time < dep_time)

# A tibble: 10,633 x 9
   origin dest  dep_delay arr_delay dep_time          sched_dep_time
  <chr>  <chr>    <dbl>    <dbl> <dtm>          <dtm>
1 EWR    BQN         9        -4 2013-01-01 19:29:00 2013-01-01 19:20:00
2 JFK    DFW        59         NA 2013-01-01 19:39:00 2013-01-01 18:40:00
3 EWR    TPA        -2         9 2013-01-01 20:58:00 2013-01-01 21:00:00
4 EWR    SJU        -6       -12 2013-01-01 21:02:00 2013-01-01 21:08:00
5 EWR    SFO        11       -14 2013-01-01 21:08:00 2013-01-01 20:57:00
6 LGA    FLL       -10        -2 2013-01-01 21:20:00 2013-01-01 21:30:00
7 EWR    MCO        41        43 2013-01-01 21:21:00 2013-01-01 20:40:00
8 JFK    LAX        -7       -24 2013-01-01 21:28:00 2013-01-01 21:35:00
9 EWR    FLL        49        28 2013-01-01 21:34:00 2013-01-01 20:45:00
10 EWR    FLL        -9       -14 2013-01-01 21:36:00 2013-01-01 21:45:00
# i 10,623 more rows
# i 3 more variables: arr_time <dtm>, sched_arr_time <dtm>, air_time <dbl>

```

These are overnight flights. We used the same date information for both the departure and the arrival times, but these flights arrived on the following day. We can fix this by adding `days(1)` to the arrival time of each overnight flight.

```

flights_dt <- flights_dt |>
  mutate(
    overnight = arr_time < dep_time,
    arr_time = arr_time + days(overnight),
    sched_arr_time = sched_arr_time + days(overnight)
  )

```

Now all of our flights obey the laws of physics.

```
flights_dt |>
  filter(arr_time < dep_time)
```

```
# A tibble: 0 x 10
# i 10 variables: origin <chr>, dest <chr>, dep_delay <dbl>, arr_delay <dbl>,
#   dep_time <dtm>, sched_dep_time <dtm>, arr_time <dtm>,
#   sched_arr_time <dtm>, air_time <dbl>, overnight <lgl>
```

Additional exercises

Calculating Elapsed Time between Two Specific Dates:

```
start_date <- ymd("1994-12-21")

end_date <- ymd("2023-11-21")

elapsed_time <- end_date - start_date

cat("Elapsed time:", as.duration(elapsed_time), "\n")
```

Elapsed time: 912556800

To enhance clarity in the output, we can represent the elapsed period in terms of days.

```
# Convert difftime to period

elapsed_period <- as.period(elapsed_time)

# Extracting components of the period

days <- day(elapsed_period)

cat("Elapsed time:", days, "days\n")
```

Elapsed time: 10562 days

Adding and subtracting months from a date:


```
future_date <- ymd("2023-05-17")

future_date_plus_3_months <- future_date + months(3)

future_date_plus_3_months
```

```
[1] "2023-08-17"
```

```
future_date_minus_2_months <- future_date - months(2)

future_date_minus_2_months
```

```
[1] "2023-03-17"
```

Extracting Components of a Date and Formatting a Date:

```
sample_date <- ymd_hms("2023-03-15 12:30:45")

formatted_date <- format(sample_date, "%A, %B %d, %Y %I:%M %p")

formatted_date
```

```
[1] "Wednesday, March 15, 2023 12:30 PM"
```

%A: Represents the full weekday name (e.g., “Sunday”, “Monday”).

%B: Represents the full month name (e.g., “January”, “February”).

%d: Represents the day of the month as a zero-padded decimal number (01, 02, ..., 31).

%Y: Represents the year with century as a decimal number.

%I: Represents the hour (01, 02, ..., 12) using the 12-hour clock.

%M: Represents the minute as a zero-padded decimal number (00, 01, ..., 59).

%p: Represents either “AM” or “PM” in uppercase.

Lubridate supports weekly periods. Here, we create a sequence of weekly periods for the next three weeks. Adding weekly periods to vector of days:

```

date_vector <- c(today(), today() + days(5), today() - weeks(2))

three_weeks = weeks(3)

future_dates <- date_vector + three_weeks

future_dates

```

```
[1] "2023-12-11" "2023-12-16" "2023-11-27"
```

Filtering dates based on a period condition :

```

date_vector <- c(today(), today() + days(5), today() - weeks(2))

filtered_dates <- date_vector[date_vector + days(3) < today()]

filtered_dates

```

```
[1] "2023-11-06"
```

In additions, we can create intervals :

```

# Working with time intervals

# Creating an interval

interval <- interval(today(), today() + hours(5))

interval

```

```
[1] 2023-11-20 UTC--2023-11-20 05:00:00 UTC
```

Then, it is possible to check if the specific date is within an interval or not :

```

# Checking if a date is within an interval

one_hour_time <- today() + hours(1)

check_within_interval <- one_hour_time %within% interval

```

```
check_within_interval
```

```
[1] TRUE
```

```
three_weeks_past = today() - weeks(3)
```

```
check_within_interval_past <- three_weeks_past %within% interval
```

```
check_within_interval_past
```

```
[1] FALSE
```